

**Київський національний університет імені Тараса Шевченка
факультет радіофізики, електроніки та комп'ютерних систем**

Лабораторна робота № 3

**Тема: «Дослідження оптимізації коду з використанням векторних
розширень CPU»**

Роботу виконав
студент 3 курсу
КІ-СА
Мургашов Гліб

Київ 2020

Хід роботи

1. Отримайте доступ на обчислювальний кластер для роботи з Intel Compiler.

[DAC746B1B36BD07A](#)

/C=UA/O=KNU/OU=People/L=FRECS/CN=Murhashov Hlib

Mon, 30 Nov 2020 20:36:52

2. Завантажте файли Intel® C++ Compiler - Using Auto-Vectorization Tutorial (<https://software.intel.com/en-us/product-code-samples?topic=20813>) на свій комп'ютер та в домашню директорію користувача обчислювального кластеру.
3. Використовуючи інструкції в readme.html ознайомтесь та виконайте Tutorial на обчислювальному кластері
Замість інструкцій в пункті "Setting the Environment Variables" завантажте оточення компілятора шляхом виконання команди: `ml icc`
Виконуйте завдання на робочих вузлах кластеру замість вхідної ноди. По-перше процесори робочих вузлів мають набагато більше розширень. По-друге виконання компіляції та запуску на вхідній ноді заважає іншим користувачам, що призведе до блокування вашого аккаунту та автоматичного незарахування лабораторної роботи.
Рекомендований варіант виконання роботи - використання інтерактивних задач в системі планування:
`[manf@plus7 ~]$ qsub -I -l nodes=1:ppn=1,walltime=00:30:00`
KNU:WN:s4 [manf ~]\$ `ml icc`

```
Last login: Thu Jun  4 09:03:33 EEST 2020 on pts/8
[tb078@plus7 ~]$ qsub -I -l nodes=1:ppn=1,walltime=00:30:00
qsub: waiting for job 2732367 to start
qsub: job 2732367 ready
```

```
autoscratch: creating directory '/mnt/work/tb078'
autoscratch: creating directory '/mnt/scratch/tb078'
KNU: :s4 [tb078 ~]$ pwd
/home/grid/testbed/tb078
```

```
KNU: :s4 [tb078 lab3]$ ls
ITAC  advisor  cluster_checker  compiler_c  compiler_f  index.html  inspector  ipp  ipsxe2019_samples_lin_20190327.tgz
KNU: :s4 [tb078 lab3]$ history
1  quit
2  exit
3  pwd
4  ml icc
5  ls -la
6  ls -la
7  mkdir lab3
8  cd lab3/
9  wget https://software.intel.com/sites/default/files/ipsxe2019_samples_lin_20190327.tgz
10 wget https://software.intel.com/sites/default/files/ipsxe2019_samples_lin_20190327.tgz
11 tar -xvf ipsxe2019_samples_lin_20190327.tgz
12 {Sls
13 ls
```

```

KNU: :s4 [tb078 vec_samples]$ cd src/
KNU: :s4 [tb078 src]$ ls
Driver.c  Multiply.c  Multiply.h
KNU: :s4 [tb078 src]$ █

```

Establishing a Performance Baseline

To set a performance baseline for the improvements that follow in this tutorial, compile your sources from the `src` directory with these compiler options:

```
icc -O1 -std=c99 Multiply.c Driver.c -o MatVector
```

Execute `MatVector` and record the execution time reported in the output. This is the baseline against which subsequent improvements will be measured.

```

KNU: :s4 [tb078 src]$ icc -O1 -std=c99 Multiply.c Driver.c -o MatVector
KNU: :s4 [tb078 src]$ ls
Driver.c  MatVector  Multiply.c  Multiply.h
KNU: :s4 [tb078 src]$ █

```

```

KNU: :s4 [tb078 src]$ ./MatVector

```

```

ROW:101 COL: 101
Execution time is 12.974 seconds
GigaFlops = 1.572524
Sum of result = 195853.999899
KNU: :s4 [tb078 src]$ █

```

A vectorization report shows what loops in your code were vectorized and explains why other loops were not vectorized. To generate a vectorization report, use the `qopt-report-phase=vec` compiler options together with `qopt-report=1` or `qopt-report=2`.

Together with `qopt-report-phase=vec`, `qopt-report=1` generates a report with the loops in your code that were vectorized while `qopt-report-phase=vec` with `qopt-report=2` generates a report with both the loops in your code that were vectorized and the reason that other loops were not vectorized.

Because vectorization is turned off with the `O1` option, the compiler does not generate a vectorization report. To generate a vectorization report, compile your project with the `O2`, `qopt-report-phase=vec`, `qopt-report=1` options:

```

KNU: :s4 [tb078 src]$ icc -std=c99 -O2 -D NOFUNCCALL -qopt-report=1 -qopt-report-phase=vec Multiply.c Driver.c -o MatVector
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location
KNU: :s4 [tb078 src]$ ls
Driver.c  Driver.optrpt  MatVector  Multiply.c  Multiply.h  Multiply.optrpt
KNU: :s4 [tb078 src]$ █

```

Recompile the program and then execute `MatVector`. Record the new execution time. The reduction in time is mostly due to auto-vectorization of the inner loop at line 145 noted in the vectorization report `matvec.optrpt`:

```

KNU: :s4 [tb078 src]$ ./MatVector

ROW:101 COL: 101
Execution time is 4.467 seconds
GigaFlops = 4.567637
Sum of result = 195853.999899
KNU: :s4 [tb078 src]$ █

```

Multiply.optrpt:

```

KNU: :s4 [tb078 src]$ cat Multiply.optrpt
Intel(R) Advisor can now assist with vectorization and show optimization
report messages with your source code.
See "https://software.intel.com/en-us/intel-advisor-xe" for details.

Begin optimization report for: matvec(int, int, double (*)[*], double *, double *)

Report from: Vector optimizations [vec]

LOOP BEGIN at Multiply.c(37,5)
remark #25460: No loop optimizations reported

LOOP BEGIN at Multiply.c(49,9)
remark #25460: No loop optimizations reported
LOOP END

LOOP BEGIN at Multiply.c(49,9)
<Remainder>
LOOP END
LOOP END
=====

```

qopt-report=2 with qopt-report-phase=vec,loop returns a list that also includes loops that were not vectorized or multi-versioned, along with the reason that the compiler did not vectorize them or multi-version the loop.

Recompile your project with the qopt-report=2 and qopt-report-phase=vec,loop options.

```

KNU: :s4 [tb078 src]$ icc -std=c99 -O2 -D NOFUNCCALL -qopt-report-phase=vec,loop -qopt-report=2 Multiply.c Driver.c -o MatVector
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location

```

```

KNU: :s4 [tb078 src]$ ./MatVector

ROW:101 COL: 101
Execution time is 4.469 seconds
GigaFlops = 4.565483
Sum of result = 195853.999899

```

The vectorization report Multiply.optrpt indicates that the loop at line 37 in Multiply.c did not vectorize because it is not the innermost loop of the loop nest. Two versions of the innermost loop at line 49 were generated, and one version was vectorized.

```

LOOP BEGIN at Multiply.c(37,5)
  remark #15541: outer loop was not auto-vectorized: consider using SIMD directive

  LOOP BEGIN at Multiply.c(49,9)
    remark #15344: loop was not vectorized: vector dependence prevents vectorization. First dependence is shown below. Use level 5 report for details
    remark #15346: vector dependence: assumed FLOW dependence between b[i] (50:13) and b[i] (50:13)
    remark #25439: unrolled with remainder by 2
  LOOP END

  LOOP BEGIN at Multiply.c(49,9)
    <Remainder>
  LOOP END
LOOP END

```

Remove the -D NOFUNCCALL to restore the call to matvec(), then add the -D NOALIAS option to the command line.

```

sum of result = 195853.999899
KNU: :s4 [tb078 src]$ icc -std=c99 -qopt-report=2 -qopt-report-phase=vec -D NOALIAS Multiply.c Driver.c -o MatVector
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location
KNU: :s4 [tb078 src]$

```

```
KNU: :s4 [tb078 src]$ ./MatVector
```

```

ROW:101 COL: 101
Execution time is 5.020 seconds
GigaFlops = 4.064306
Sum of result = 195853.999899
-----

```

This conditional compilation replaces the loop in the main program with a function call. Execute MatVector and record the execution time reported in the output. Multiply.optrpt now shows:

```

KNU: :s4 [tb078 src]$ cat Multiply.optrpt
Intel(R) Advisor can now assist with vectorization and show optimization
report messages with your source code.
See "https://software.intel.com/en-us/intel-advisor-xe" for details.

```

```
Begin optimization report for: matvec(int, int, double (*)[*], double *__restrict__, double *)
```

```
Report from: Vector optimizations [vec]
```

```

LOOP BEGIN at Multiply.c(37,5)
  remark #15542: loop was not vectorized: inner loop was already vectorized

  LOOP BEGIN at Multiply.c(49,9)
    <Peeled loop for vectorization>
  LOOP END

  LOOP BEGIN at Multiply.c(49,9)
    remark #15300: LOOP WAS VECTORIZED
  LOOP END

  LOOP BEGIN at Multiply.c(49,9)
    <Alternate Alignment Vectorized Loop>
  LOOP END

  LOOP BEGIN at Multiply.c(49,9)
    <Remainder loop for vectorization>
  LOOP END
LOOP END
=====

```

The vectorizer can generate faster code when operating on aligned data. In this activity you will improve performance by aligning the arrays a, b, and x in Driver.c on a 16-byte boundary so that the vectorizer can use aligned load instructions for all arrays rather than the slower unaligned load instructions and can avoid runtime tests of alignment. Using the ALIGNED macro will modify the declarations of a, b, and x in Driver.c using the aligned attribute keyword

Recompile the program after adding the ALIGNED macro to ensure consistently aligned data. Use -qopt-report=4 to see the change in aligned references.

```
=====  
KNU: :s4 [tb078 src]$ icc -std=c99 -qopt-report=4 -qopt-report-phase=vec -D NOALIAS -D ALIGNED Multiply.c Driver.c -o MatVector  
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location
```

```
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location  
KNU: :s4 [tb078 src]$ ./MatVector
```

```
ROW:101 COL: 102  
Execution time is 4.618 seconds  
GigaFlops = 4.418266  
Sum of result = 195853.999899  
KNU: :s4 [tb078 src]$
```

Multiply.optrpt after adding -D ALIGNED shows:

```
Sum of result = 195853.999899  
KNU: :s4 [tb078 src]$ cat Multiply.optrpt  
Intel(R) Advisor can now assist with vectorization and show optimization  
report messages with your source code.  
See "https://software.intel.com/en-us/intel-advisor-xe" for details.  
  
Intel(R) C Intel(R) 64 Compiler for applications running on Intel(R) 64, Version 18.0.5.274 Build 20180823  
  
Compiler options: -std=c99 -qopt-report=4 -qopt-report-phase=vec -D NOALIAS -D ALIGNED -o MatVector  
  
Begin optimization report for: matvec(int, int, double (*)[*], double *__restrict__, double *)  
  
Report from: Vector optimizations [vec]  
  
LOOP BEGIN at Multiply.c(37,5)  
remark #15542: loop was not vectorized: inner loop was already vectorized  
  
LOOP BEGIN at Multiply.c(49,9)  
remark #15388: vectorization support: reference a[i][j] has aligned access [ Multiply.c(50,21) ]  
remark #15388: vectorization support: reference x[j] has aligned access [ Multiply.c(50,31) ]  
remark #15305: vectorization support: vector length 2  
remark #15399: vectorization support: unroll factor set to 4  
remark #15309: vectorization support: normalized vectorization overhead 0.594  
remark #15300: LOOP WAS VECTORIZED  
remark #15448: unmasked aligned unit stride loads: 2  
remark #15475: --- begin vector cost summary ---  
remark #15476: scalar cost: 10  
remark #15477: vector cost: 4.000  
remark #15478: estimated potential speedup: 2.410  
remark #15488: --- end vector cost summary ---  
LOOP END  
  
LOOP BEGIN at Multiply.c(49,9)  
<Remainder loop for vectorization>  
remark #15388: vectorization support: reference a[i][j] has aligned access [ Multiply.c(50,21) ]  
remark #15388: vectorization support: reference x[j] has aligned access [ Multiply.c(50,31) ]  
remark #15335: remainder loop was not vectorized: vectorization possible but seems inefficient. Use vector always directive or -vec-threshold0 to override  
remark #15305: vectorization support: vector length 2  
remark #15309: vectorization support: normalized vectorization overhead 2.417  
LOOP END  
LOOP END  
=====
```

The compiler may be able to perform additional optimizations if it is able to optimize across source line boundaries. These may include, but are not limited to, function inlining. This is enabled with the -ipo option.

Recompile the program using the -ipo option to enable interprocedural optimization.

```
driver.c driver.optrpt MatVector multiply.c multiply.in multiply.optrpt ipo_out.optrpt
KNU: :s4 [tb078 src]$ icc -std=c99 -qopt-report=2 -qopt-report-phase=vec -D NOALIAS -D ALIGNED -ipo Multiply.c Driver.c -o MatVector
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location
```

```
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location
KNU: :s4 [tb078 src]$ ./MatVector
```

```
ROW:101 COL: 102
Execution time is 4.329 seconds
GigaFlops = 4.712376
Sum of result = 195853.999899
KNU: :s4 [tb078 src]$
```

Note that the vectorization messages now appear at the point of inlining in Driver.c (line 150) and this is found in the file ipo_out.optrpt.

```
KNU: :s4 [tb078 src]$ cat ipo_out.optrpt
Intel(R) Advisor can now assist with vectorization and show optimization
report messages with your source code.
See "https://software.intel.com/en-us/intel-advisor-xe" for details.

Begin optimization report for: main()

Report from: Vector optimizations [vec]

LOOP BEGIN at Driver.c(152,16)
  remark #15542: loop was not vectorized: inner loop was already vectorized

LOOP BEGIN at Multiply.c(37,5) inlined into Driver.c(150,9)
  remark #15542: loop was not vectorized: inner loop was already vectorized

  LOOP BEGIN at Multiply.c(49,9) inlined into Driver.c(150,9)
    remark #15300: LOOP WAS VECTORIZED
  LOOP END

  LOOP BEGIN at Multiply.c(49,9) inlined into Driver.c(150,9)
    <Remainder loop for vectorization>
    remark #15335: remainder loop was not vectorized: vectorization possible but seems inefficient. Use vector always directive or -vec-threshold0 to override
  LOOP END
LOOP END

LOOP BEGIN at Driver.c(74,5) inlined into Driver.c(159,5)
  remark #15300: LOOP WAS VECTORIZED
LOOP END

LOOP BEGIN at Driver.c(74,5) inlined into Driver.c(159,5)
  <Remainder loop for vectorization>
  LOOP END
=====
```

```
=====
Begin optimization report for: init_matrix(int, int, double, double (*)(102))
    Report from: Vector optimizations [vec]
```

```
LOOP BEGIN at Driver.c(47,5)
    remark #15542: loop was not vectorized: inner loop was already vectorized

    LOOP BEGIN at Driver.c(48,9)
        remark #15300: LOOP WAS VECTORIZED
    LOOP END

    LOOP BEGIN at Driver.c(48,9)
        <Remainder loop for vectorization>
    LOOP END
LOOP END

LOOP BEGIN at Driver.c(53,9)
    remark #15300: LOOP WAS VECTORIZED
LOOP END

LOOP BEGIN at Driver.c(53,9)
    <Remainder loop for vectorization>
LOOP END
=====
```

```
Begin optimization report for: init_array(int, double, double *)
    Report from: Vector optimizations [vec]
```

```
LOOP BEGIN at Driver.c(62,5)
    remark #15300: LOOP WAS VECTORIZED
LOOP END

LOOP BEGIN at Driver.c(62,5)
    <Remainder loop for vectorization>
LOOP END
=====
```


4. Оберіть будь-яку неінтерактивну консольну програму мовою C/C++ (унікальну в межах групи, в гуглі більше ніж 50 програм)

Взята одна із програм з минулорічної практики, яка виконує дії з матрицями.

https://github.com/HleBASS/CS_labs/blob/master/Lab3/myprog.cpp

Напишіть сценарій, що:

Компілює програму з різними оптимізаціями (-O) та виміряйте час її роботи. Якщо час досить малий - вимірюйте час роботи 1000 (чи 1000000) запусків алгоритму в циклі. Час роботи можна виміряти утилітою time.

Отримує перелік всіх розширень процесору що підтримуються
Для кожного розширення компілює Intel-компілятором окремий варіант оптимізованого коду (наприклад -x SSE2)

Вимірює час виконання кожного варіанта оптимізованої програми
Запустіть задачу в планувальник обчислювального кластеру 5 разів (для статистики на різних нодах)

```
[manf@plus7 ~]$ qsub -N MyJob -l nodes=1:ppn=1,walltime=00:30:00  
script.sh
```


Побудуйте графіки залежності часу від різних варіантів компіляції.

Був написаний скрипт, який замірює час виконання циклу з 100 запусків виконання оптимізованої програми.

https://github.com/HleBASS/CS_labs/blob/master/Lab3/script.sh

Час виконання програми скомпільованої з різними прапорцями.


Номер запуску \ прапорець	O0	O1	O2	O3	Os	Ofast
1	3.059	3.98	3.033	1.228	3.107	0.485
2	3.054	3.102	3.049	1.182	3.073	0.438
3	3.079	3.114	3.078	1.199	3.085	0.454
4	3.073	3.99	3.065	1.198	3.06	0.492
5	3.122	3.143	3.107	1.241	3.128	0.499
average	3.077	3.466	3.066	1.209	3.09	0.474



Результат з -Ofast прапорцем і різними розширеннями:


Номер запуску \ розширення	sse	sse2	ssse3	avx
1	0.452	0.452	0.452	0.453
2	0.427	0.323	0.437	0.438
3	0.422	0.436	0.434	0.434
4	0.455	0.457	0.453	0.456
5	0.457	0.456	0.459	0.456
average	0.443	0.425	0.447	0.447

В даному випадку за допомогою розширень вдалося зекономити ще ~43% часу виконання програми.



Result with -O1 flag and different extensions:

Номер запуску \ розширення	sse	sse2	ssse3	avx
1	2.218	2.184	2.183	2.183
2	2.143	2.149	2.129	2.155
3	2.154	2.152	2.143	2.149
4	2.193	2.199	2.188	2.191
5	2.197	2.193	2.197	2.194
average	2.181	2.176	2.168	2.174



Висновок: В даній лабораторній роботі було ознайомлено з автоматичною векторизацією в паралельних обчисленнях. Було проведено оптимізацію коду з використанням векторних розширень CPU компілятора Intel на прикладі програми, яка виконує заповнення масивів.